

# CS61B, 2019

---

## Lecture 19: Hashing

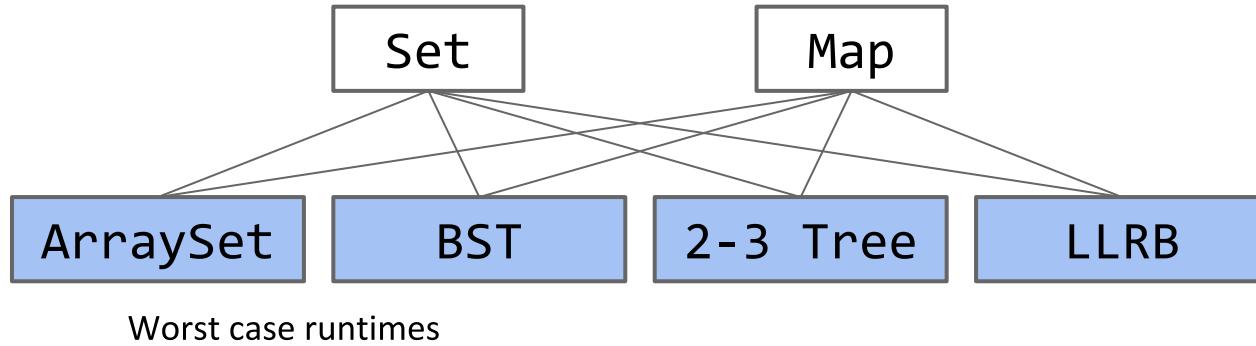
- Set Implementations, DataIndexedIntegerSet
- Integer Representations of Strings, Integer Overflow
- Hash Tables and Handling Collisions
- Hash Table Performance and Resizing
- Hash Tables in Java



# Data Indexed Arrays

# Sets

We've now seen several implementations of the Set (or Map) ADT.



	contains(x)	add(x)	Notes
ArraySet	$\Theta(N)$	$\Theta(N)$	
BST	$\Theta(N)$	$\Theta(N)$	Random trees are $\Theta(\log N)$ .
2-3 Tree	$\Theta(\log N)$	$\Theta(\log N)$	Beautiful idea. Very hard to implement.
LLRB	$\Theta(\log N)$	$\Theta(\log N)$	Maintains bijection with 2-3 tree. Hard to implement.

# Limits of Search Tree Based Sets

---

Our search tree based sets require items to be comparable.

- Need to be able to ask “is  $X < Y$ ?” Not true of all types.
- Could we somehow avoid the need for objects to be comparable?

Our search tree sets have excellent performance, but could maybe be better?

- $\Theta(\log N)$  is amazing. 1 billion items is still only height  $\sim 30$ .
- Could we somehow do better than  $\Theta(\log N)$ ?

Today we'll see the answer to both of the questions above is yes.

# Using Data as an Index

One extreme approach: Create an array of booleans indexed by data!

- Initially all values are false.
- When an item is added, set appropriate index to true.

```
DataIndexedIntegerSet diis;  
diis = new DataIndexedIntegerSet();  
diis.add(0);
```

Set containing nothing

F	0
F	1
F	2
F	3
F	4
F	5
F	6
F	7
F	8
F	9
F	10
F	11
F	12
F	13
F	14
F	15

...

# Using Data as an Index

One extreme approach: Create an array of booleans indexed by data!

- Initially all values are false.
- When an item is added, set appropriate index to true.

```
DataIndexedIntegerSet diis;  
diis = new DataIndexedIntegerSet();  
diis.add(0);
```

Set containing 0

T	0
F	1
F	2
F	3
F	4
F	5
F	6
F	7
F	8
F	9
F	10
F	11
F	12
F	13
F	14
F	15

...

# Using Data as an Index

One extreme approach: Create an array of booleans indexed by data!

- Initially all values are false.
- When an item is added, set appropriate index to true.

```
DataIndexedIntegerSet diis;  
diis = new DataIndexedIntegerSet();  
diis.add(0);  
diis.add(5);
```

Set containing 0, 5

T	0
F	1
F	2
F	3
F	4
T	5
F	6
F	7
F	8
F	9
F	10
F	11
F	12
F	13
F	14
F	15

...

# Using Data as an Index

One extreme approach: Create an array of booleans indexed by data!

- Initially all values are false.
- When an item is added, set appropriate index to true.

```
DataIndexedIntegerSet diis;  
diis = new DataIndexedIntegerSet();  
diis.add(0);  
diis.add(5);  
diis.add(10);
```

Set containing 0, 5, 10

T	0
F	1
F	2
F	3
F	4
T	5
F	6
F	7
F	8
F	9
T	10
F	11
F	12
F	13
F	14
F	15

...

# Using Data as an Index

One extreme approach: Create an array of booleans indexed by data!

- Initially all values are false.
- When an item is added, set appropriate index to true.

```
DataIndexedIntegerSet diis;  
diis = new DataIndexedIntegerSet();  
diis.add(0);  
diis.add(5);  
diis.add(10);  
diis.add(11);
```

Set containing 0, 5, 10, 11

T	0
F	1
F	2
F	3
F	4
T	5
F	6
F	7
F	8
F	9
T	10
T	11
F	12
F	13
F	14
F	15

...

# DataIndexedIntegerSet Implementation

```
public class DataIndexedIntegerSet {  
    private boolean[] present;  
  
    public DataIndexedIntegerSet() {  
        present = new boolean[2000000000];  
    }  
  
    public void add(int i) {  
        present[i] = true;  
    }  
  
    public boolean contains(int i) {  
        return present[i];  
    }  
}
```

T	0
F	1
F	2
F	3
F	4
T	5
F	6
F	7
F	8
F	9
T	10
T	11
F	12
F	13
F	14
F	15

...

# DataIndexedIntegerSet Implementation

```
public class DataIndexedIntegerSet {  
    private boolean[] present;  
  
    public DataIndexedIntegerSet() {  
        present = new boolean[2000000000];  
    }  
  
    public void add(int i) {  
        present[i] = true;  
    }  
  
    public boolean contains(int i) {  
        return present[i];  
    }  
}
```

	contains(x)	add(x)
ArraySet	$\Theta(N)$	$\Theta(N)$
BST	$\Theta(N)$	$\Theta(N)$
2-3 Tree	$\Theta(\log N)$	$\Theta(\log N)$
LLRB	$\Theta(\log N)$	$\Theta(\log N)$
DataIndexedArray	$\Theta(1)$	$\Theta(1)$

# Using Data as an Index

One extreme approach: Create an array of booleans indexed by data!

- Initially all values are false.
- When an item is added, set appropriate index to true.

Downsides of this approach (that we will try to address):

- Extremely wasteful of memory. To support checking presence of all positive integers, we need > 2 billion booleans.
- Need some way to generalize beyond integers.

```
DataIndexedIntegerSet diis;  
diis = new DataIndexedIntegerSet();  
diis.add(0);  
diis.add(5);  
diis.add(10);  
diis.add(11);
```

Set containing 0, 5, 10, 11

T	0
F	1
F	2
F	3
F	4
T	5
F	6
F	7
F	8
F	9
T	10
T	11
F	12
F	13
F	14
F	15

...

# DataIndexedEnglishWordSet



# Generalizing the DataIndexedIntegerSet Idea

Ideally, we want a data indexed set that can store arbitrary types.

```
DataIndexedSet<String> dis =  
    new DataIndexedSet<>();  
  
dis.add("cat");  
dis.add("bee");  
dis.add("dog");
```

Where do cat, bee,  
and dog go???

F	0
F	1
F	2
F	3
F	4
F	5
F	6
...	

But previous idea only supports integers!

- Let's talk about storing Strings.
- Will get to generic types later.



# Generalizing the DataIndexedIntegerSet Idea

Suppose we want to add("cat")

The key question:

- What is the **cat**th element of a list?
- One idea: Use the first letter of the word as an index.
  - **a** = 1, **b** = 2, **c** = 3, ..., **z** = 26

0	F	
1	F	a
2	F	b
3	T	c
4	F	d
...		
25	F	y
26	F	z

What's wrong with this approach?

- Other words start with **c**.
  - contains("chupacabra") : true
- Can't store "=98yaе98fwуawef"

"chupacabra" **collides** with "cat"



# Avoiding Collisions

Use all digits by multiplying each by a power of 27.

- $a = 1, b = 2, c = 3, \dots, z = 26$
- Thus the index of “cat” is  $(3 \times 27^2) + (1 \times 27^1) + (20 \times 27^0)$   
= 2234.



Why this specific pattern?

- Let's review how numbers are represented in decimal.

0	F	
1	F	a
2	F	b
3	F	c
4	F	d
...		

25	F	y
26	F	z

...

2233	F	cas
2234	T	cat
2235	F	cau

...

# The Decimal Number System vs. Our System for Strings

---

In the decimal number system, we have 10 digits: 0, 1, 2, 3, 4, 5, 6, 7, 8, 9

Want numbers larger than 9? Use a sequence of digits.

Example: 7091 in base 10

- $7091_{10} = (7 \times 10^3) + (0 \times 10^2) + (9 \times 10^1) + (1 \times 10^0)$

Our system for strings is almost the same, but with letters.

## Test Your Understanding

---

Convert the word “bee” into a number by using our “powers of 27” strategy.

Reminder:  $\text{cat}_{27} = (3 \times 27^2) + (1 \times 27^1) + (20 \times 27^0) = 2234_{10}$

Hint: ‘b’ is letter 2, and ‘e’ is letter 5.

## Test Your Understanding

---

Convert the word “bee” into a number by using our “powers of 27” strategy.

Reminder:  $\text{cat}_{27} = (3 \times 27^2) + (1 \times 27^1) + (20 \times 27^0) = 2234_{10}$

Hint: ‘b’ is letter 2, and ‘e’ is letter 5.

- $\text{bee}_{27} = (2 \times 27^2) + (5 \times 27^1) + (5 \times 27^0) = 1598_{10}$

## Uniqueness

---

- $\text{cat}_{27} = (3 \times 27^2) + (1 \times 27^1) + (20 \times 27^0) = 2234_{10}$
- $\text{bee}_{27} = (2 \times 27^2) + (5 \times 27^1) + (5 \times 27^0) = 1598_{10}$

As long as we pick a base  $\geq 26$ , this algorithm is guaranteed to give each lowercase English word a unique number!

- Using base 27, no other words will get the number 1598.

In other words: Guaranteed that we will never have a collision.

## Implementing englishToInt (optional)

---

Optional exercise: Try to write a function `englishToInt` that can convert English strings to integers by adding characters scaled by powers of 27.

Examples:

- a: 1
- z: 26
- aa: 28
- bee: 1598
- cat: 2234
- dog: 3328
- potato: 237,949,071

# Implementing englishToInt (optional) (solution)

```
/** Converts ith character of String to a letter number.  
 * e.g. 'a' -> 1, 'b' -> 2, 'z' -> 26 */  
public static int letterNum(String s, int i) {  
    int ithChar = s.charAt(i);  
    if ((ithChar < 'a') || (ithChar > 'z'))  
        { throw new IllegalArgumentException(); }  
    return ithChar - 'a' + 1;  
}  
  
public static int englishToInt(String s) {  
    int intRep = 0;  
    for (int i = 0; i < s.length(); i += 1) {  
        intRep = intRep * 27;  
        intRep = intRep + letterNum(s, i);  
    }  
    return intRep;  
}
```

# DataIndexedEnglishWordSet Implementation

```
public class DataIndexedEnglishWordSet {  
    private boolean[] present;  
  
    public DataIndexedEnglishWordSet() {  
        present = new boolean[2000000000];  
    }  
  
    public void add(String s) {  
        present[englishToInt(s)] = true;  
    }  
  
    public boolean contains(int i) {  
        return present[englishToInt(s)];  
    }  
}
```

0	F	
1	F	a
2	F	b
3	F	c
4	F	d
...		
25	F	y
26	F	z
...		
2233	F	cas
2234	T	cat
2235	F	cau
...		

Set containing “cat”

# DataIndexedStringSet



# DataIndexedStringSet

---

Using only lowercase English characters is too restrictive.

- What if we want to store strings like “2pac” or “eGg!”?
- To understand what value we need to use for our base, let’s discuss briefly discuss the ASCII standard.

# ASCII Characters

The most basic character set used by most computers is ASCII format.

- Each possible character is assigned a value between 0 and 127.
- Characters 33 - 126 are “printable”, and are shown below.
- For example, **char c = 'D'** is equivalent to **char c = 68**.

33	!	49	1	65	A	81	Q	97	a	113	q
34	"	50	2	66	B	82	R	98	b	114	r
35	#	51	3	67	C	83	S	99	c	115	s
36	\$	52	4	68	D	84	T	100	d	116	t
37	%	53	5	69	E	85	U	101	e	117	u
38	&	54	6	70	F	86	V	102	f	118	v
39	'	55	7	71	G	87	W	103	g	119	w
40	(	56	8	72	H	88	X	104	h	120	x
41	)	57	9	73	I	89	Y	105	i	121	y
42	*	58	:	74	J	90	Z	106	j	122	z
43	+	59	;	75	K	91	[	107	k	123	{
44	,	60	<	76	L	92	\	108	l	124	
45	-	61	=	77	M	93	]	109	m	125	}
46	.	62	>	78	N	94	^	110	n	126	~
47	/	63	?	79	O	95	-`	111	o		
48	0	64	@	80	P	96		112	p		

biggest value is 126

# DataIndexedStringSet

---

Using only lowercase English characters is too restrictive.

- What if we want to store strings like “2pac” or “eGg!”?
- Maximum possible value for english-only text including punctuation is 126, so let’s use 126 as our base in order to ensure unique values for all possible strings.

Examples:

- $\text{bee}_{126} = (98 \times 126^2) + (101 \times 126^1) + (101 \times 126^0) = 1,568,675$
- $\text{2pac}_{126} = (50 \times 126^3) + (112 \times 126^2) + (97 \times 126^1) + (99 \times 126^0) = 101,809,233$
- $\text{eGg!}_{126} = (98 \times 126^3) + (71 \times 126^2) + (98 \times 126^1) + (33 \times 126^0) = 203,178,213$

# Implementing asciiToInt

The corresponding integer conversion function is actually even simpler than englishToInt! Using the raw character value means we avoid the need for a helper method.

```
public static int asciiToInt(String s) {  
    int intRep = 0;  
    for (int i = 0; i < s.length(); i += 1) {  
        intRep = intRep * 126;  
        intRep = intRep + s.charAt(i);  
    }  
    return intRep;  
}
```

What if we want to use characters beyond ASCII?

# Going Beyond ASCII

---

chars in Java also support character sets for other languages and symbols.

- `char c = '☂'` is equivalent to `char c = 9730`.
- `char c = '鳌'` is equivalent to `char c = 40140`.
- `char c = 'Ѐ'` is equivalent to `char c = 54812`.
- This encoding is known as Unicode. Table is too big to list.



## Example: Computing Unique Representations of Chinese

The largest possible value for Chinese characters is 40,959\*, so we'd need to use this as our base if we want to have a unique representation for all possible strings of Chinese characters.

Example:

- 守门员<sub>40959</sub> = (23432 x 40959<sup>2</sup>) + (38376 x 40959<sup>1</sup>) + (21592 x 40959<sup>0</sup>) = 39,312,024,869,368

\*If you're curious, the last character is: 犬

39,3120,2486,9367	...	F	守门员
39,3120,2486,9368		T	守门员
39,3120,2486,9369		F	守门员

# **Integer Overflow and Hash Codes**



# Major Problem: Integer Overflow

---

In Java, the largest possible integer is 2,147,483,647.

- If you go over this limit, you overflow, starting back over at the smallest integer, which is -2,147,483,648.
- In other words, the next number after 2,147,483,647 is -2,147,483,648.

```
int x = 2147483647;
System.out.println(x);
System.out.println(x + 1);
```

```
jug ~/Dropbox/61b/lec/hashing
$ javac BiggestPlusOne.java
$ java BiggestPlusOne
2147483647
-2147483648
```

## Consequence of Overflow: Collisions

---

Because Java has a maximum integer, we won't get the numbers we expect!

- With base 126, we will run into overflow even for short strings.
  - Example:  $\text{omens}_{126} = 28,196,917,171$ , which is much greater than the maximum integer!
  - `asciiToInt('omens')` will give us -1,867,853,901 instead.

# Consequence of Overflow: Collisions

Because Java has a maximum integer, we won't get the numbers we expect!

- With base 126, we will run into overflow even for short strings.
  - Example:  $\text{omens}_{126} = 28,196,917,171$ , which is much greater than the maximum integer!
  - `asciiToInt('omens')` will give us -1,867,853,901 instead.

Overflow can result in **collisions**, causing incorrect answers due.

```
public void moo() {  
    DataIndexedStringSet disi = new DataIndexedStringSet();  
    disi.add("melt banana");  
    disi.contains("subterrestrial anticosmetic"); ← returns true!  
    //asciiToInt for these strings is 839099497  
}
```

# Hash Codes and the Pigeonhole Principle

---

The official term for the number we're computing is "hash code".

- Via [Wolfram Alpha](#): a hash code "projects a value from a set with many (or even an infinite number of) members to a value from a set with a fixed number of (fewer) members."
- Here, our target set is the set of Java integers, which is of size 4294967296.

# Hash Codes and the Pigeonhole Principle

---

The official term for the number we're computing is "hash code".

- Via [Wolfram Alpha](#): a hash code "projects a value from a set with many (or even an infinite number of) members to a value from a set with a fixed number of (fewer) members."
- Here, our target set is the set of Java integers, which is of size 4294967296.

[Pigeonhole principle](#) tells us that if there are more than 4294967296 possible items, multiple items will share the same hash code.

- There are more than 4294967296 planets.
  - Each has mass, xPos, yPos, xVel, yVel, imgName.
- There are more than 4294967296 strings.
  - "one", "two", ... "nineteen quadrillion", ...



**Bottom line: Collisions are inevitable.**

# Two Fundamental Challenges

---

Two fundamental challenges:

- How do we resolve hashCode collisions (“melt banana” vs. “subterranean anticosmetic”)?
  - We’ll call this ***collision handling***.
- How do we compute a hash code for arbitrary objects?
  - We’ll call this ***computing a hashCode***.
  - Example: Our hashCode for “melt banana” was 839099497.
  - For Strings, this was relatively straightforward (treat as a base 27 or base 126 or base 40959 number).

# Hash Tables: Handling Collisions



# Resolving Ambiguity

Pigeonhole principle tells us that collisions are inevitable due to integer overflow.

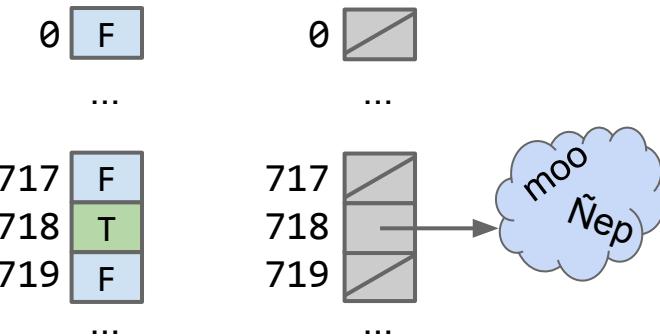
- Example: hash code for “moo” and “Ñep” might both be 718.

Suppose N items have the same numerical representation h:

- Instead of storing true in position h, store a “bucket” of these N items at position h.

How to implement a “bucket”?

- Conceptually simplest way: LinkedList.
- Could also use ArrayLists.
- Could also use an ArraySet.
- Will see it doesn't really matter what you do.



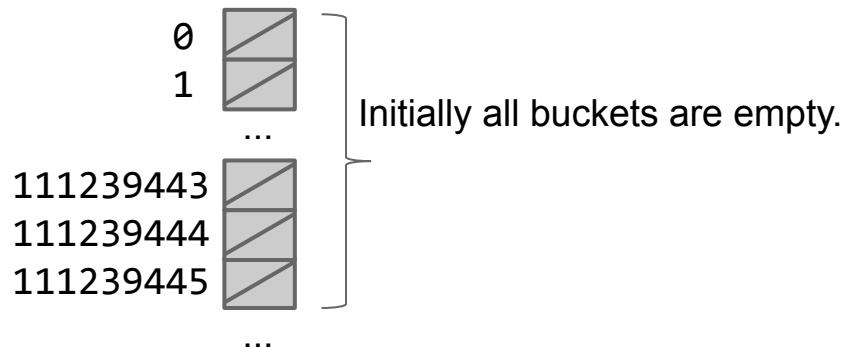
# The Separate Chaining Data Indexed Array

Each bucket in our array is initially empty. When an item  $x$  gets added at index  $h$ :

- If bucket  $h$  is empty, we create a new list containing  $x$  and store it at index  $h$ .
- If bucket  $h$  is already a list, we add  $x$  to this list if it is not already present.

We might call this a “separate chaining data indexed array”.

- Bucket # $h$  is a “separate chain” of all items that have hash code  $h$ .



# The Separate Chaining Data Indexed Array

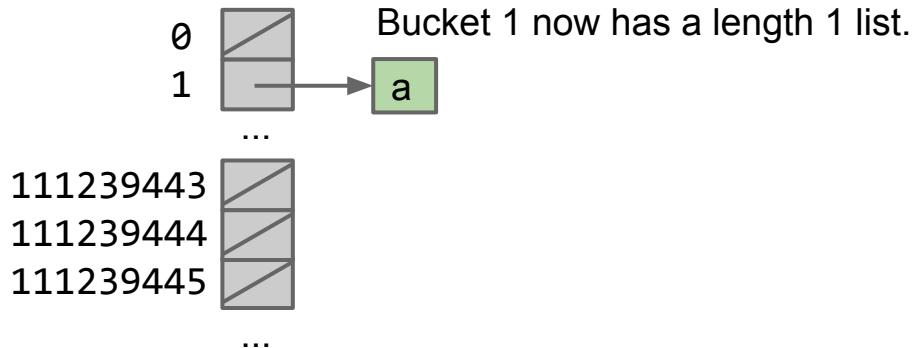
Each bucket in our array is initially empty. When an item  $x$  gets added at index  $h$ :

- If bucket  $h$  is empty, we create a new list containing  $x$  and store it at index  $h$ .
- If bucket  $h$  is already a list, we add  $x$  to this list if it is not already present.

We might call this a “separate chaining data indexed array”.

- Bucket # $h$  is a “separate chain” of all items that have hash code  $h$ .

add(“a”)



# The Separate Chaining Data Indexed Array

Each bucket in our array is initially empty. When an item  $x$  gets added at index  $h$ :

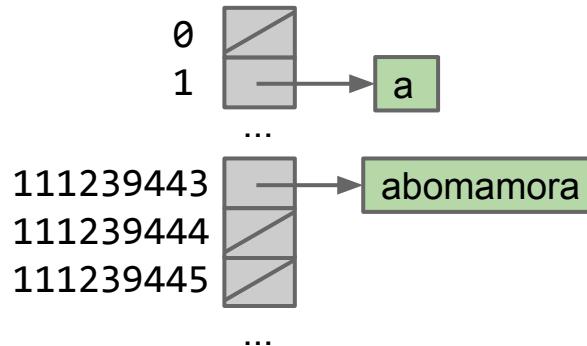
- If bucket  $h$  is empty, we create a new list containing  $x$  and store it at index  $h$ .
- If bucket  $h$  is already a list, we add  $x$  to this list if it is not already present.

We might call this a “separate chaining data indexed array”.

- Bucket # $h$  is a “separate chain” of all items that have hash code  $h$ .

`add("a")`

`add("abomamora")`



# The Separate Chaining Data Indexed Array

Each bucket in our array is initially empty. When an item  $x$  gets added at index  $h$ :

- If bucket  $h$  is empty, we create a new list containing  $x$  and store it at index  $h$ .
- If bucket  $h$  is already a list, we add  $x$  to this list if it is not already present.

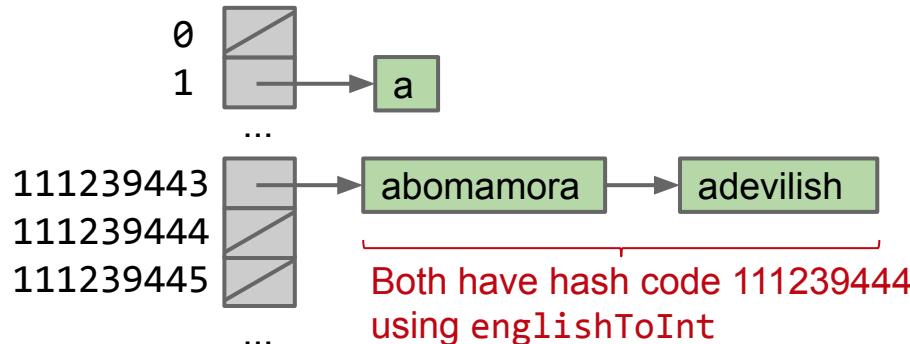
We might call this a “separate chaining data indexed array”.

- Bucket # $h$  is a “separate chain” of all items that have hash code  $h$ .

`add("a")`

`add("abomamora")`

`add("adevilish")`



# The Separate Chaining Data Indexed Array

Each bucket in our array is initially empty. When an item  $x$  gets added at index  $h$ :

- If bucket  $h$  is empty, we create a new list containing  $x$  and store it at index  $h$ .
- If bucket  $h$  is already a list, we add  $x$  to this list if it is not already present.

We might call this a “separate chaining data indexed array”.

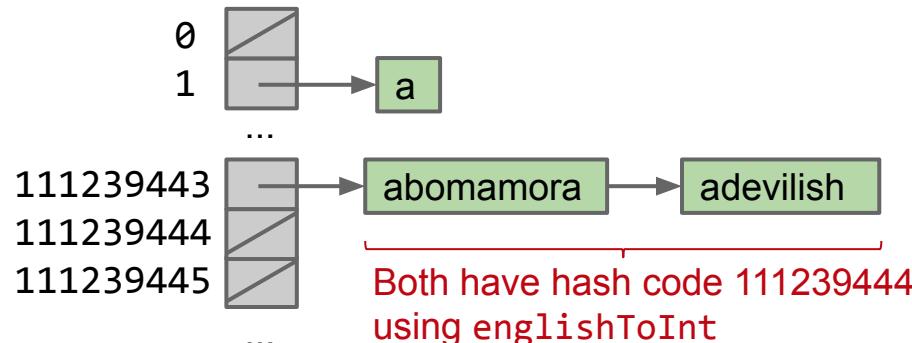
- Bucket # $h$  is a “separate chain” of all items that have hash code  $h$ .

`add("a")`

`add("abomamora")`

`add("adevilish")`

`add("abomamora")`



# The Separate Chaining Data Indexed Array

Each bucket in our array is initially empty. When an item  $x$  gets added at index  $h$ :

- If bucket  $h$  is empty, we create a new list containing  $x$  and store it at index  $h$ .
- If bucket  $h$  is already a list, we add  $x$  to this list if it is not already present.

We might call this a “separate chaining data indexed array”.

- Bucket # $h$  is a “separate chain” of all items that have hash code  $h$ .

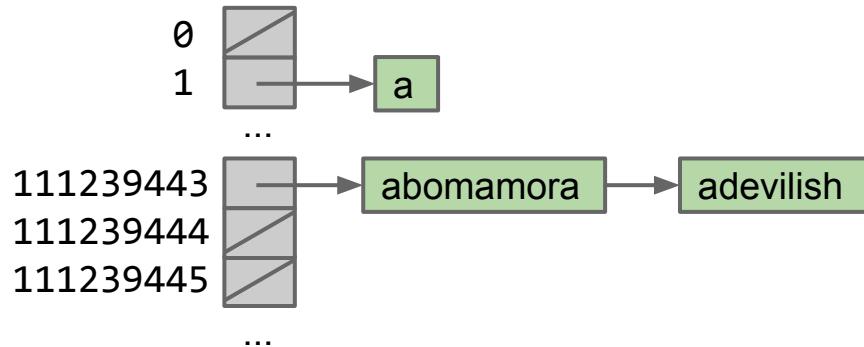
`add("a")`

`add("abomamora")`

`add("adevilish")`

`add("abomamora")`

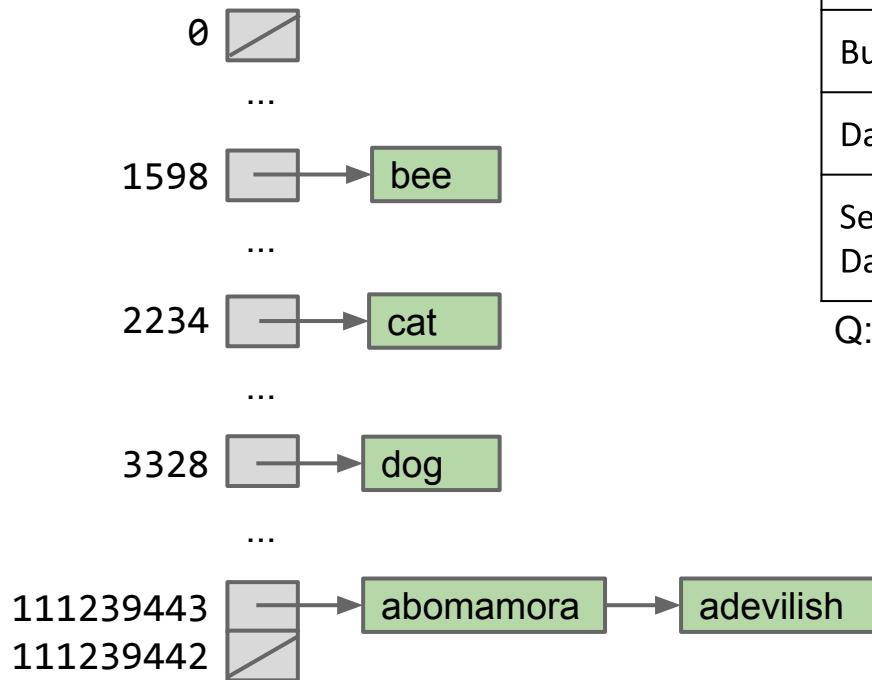
`contains("adevilish")`



- Look at all items in bucket 111239443 to see if “adevilish” is present.

# Separate Chaining Performance

Observation: Worst case runtime will be proportional to length of longest list.

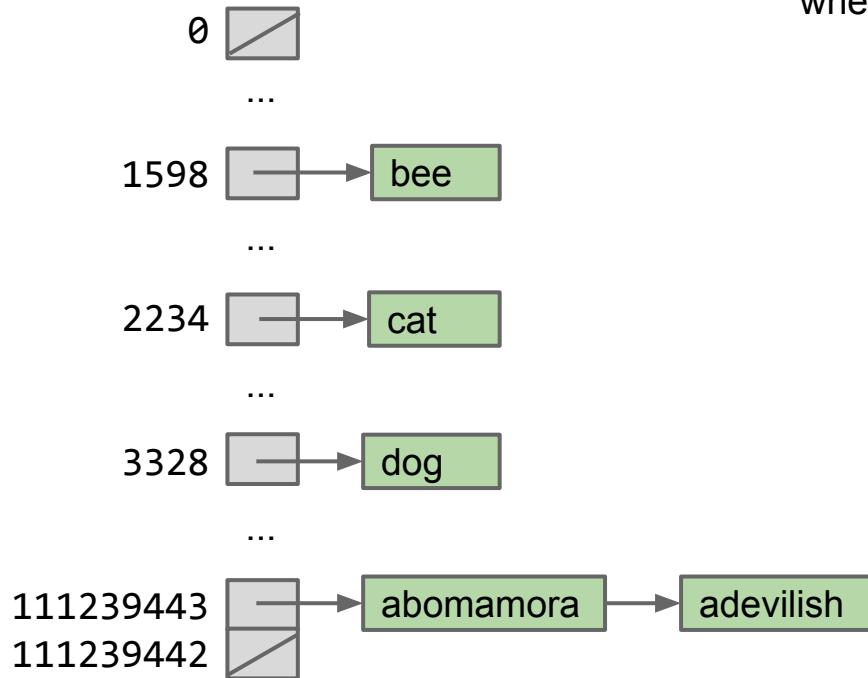


Worst case time	contains(x)	add(x)
Bushy BSTs	$\Theta(\log N)$	$\Theta(\log N)$
DataIndexedArray	$\Theta(1)$	$\Theta(1)$
Separate Chaining Data Indexed Array	$\Theta(Q)$	$\Theta(Q)$

Q: Length of longest list      Why Q and not 1?

# Saving Memory Using Separate Chaining

Observation: We don't really need billions of buckets.



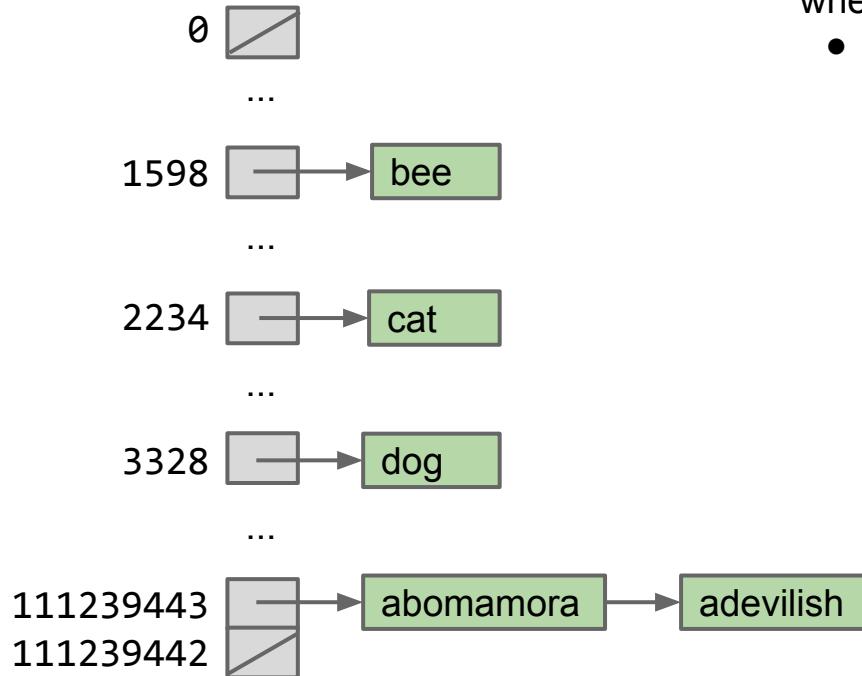
Q: If we use the 10 buckets on the right, where should our five items go?



# Saving Memory Using Separate Chaining and Modulus

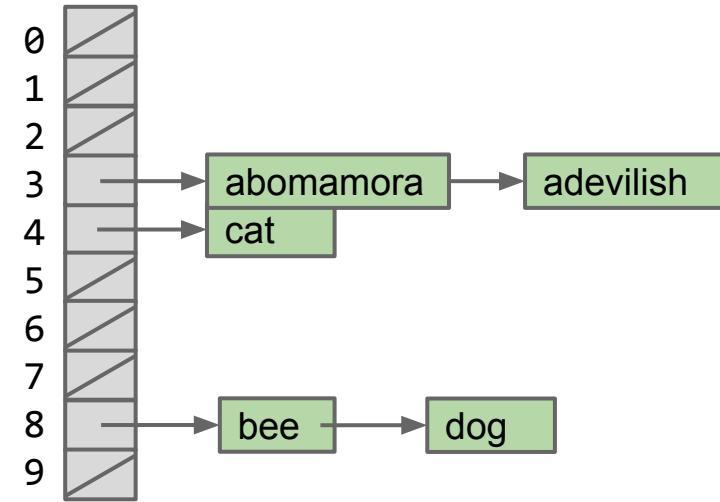
Observation: Can use modulus of hashCode to reduce bucket count.

- Downside: Lists will be longer.



Q: If we use the 10 buckets on the right, where should our five items go?

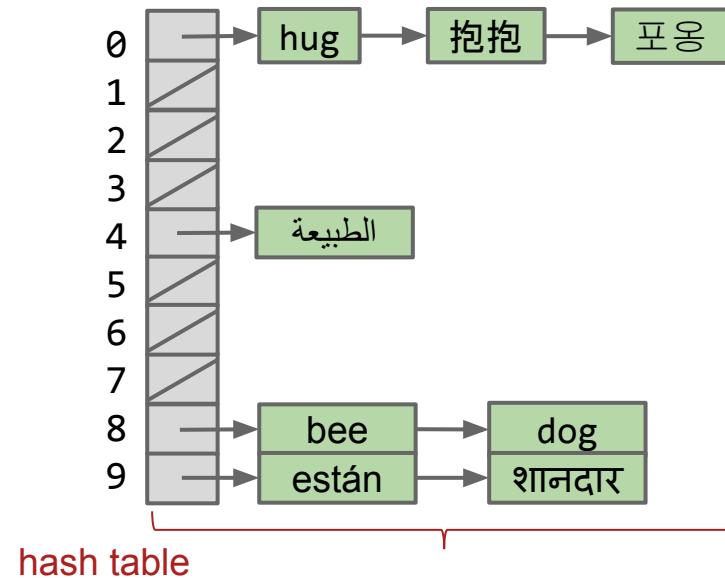
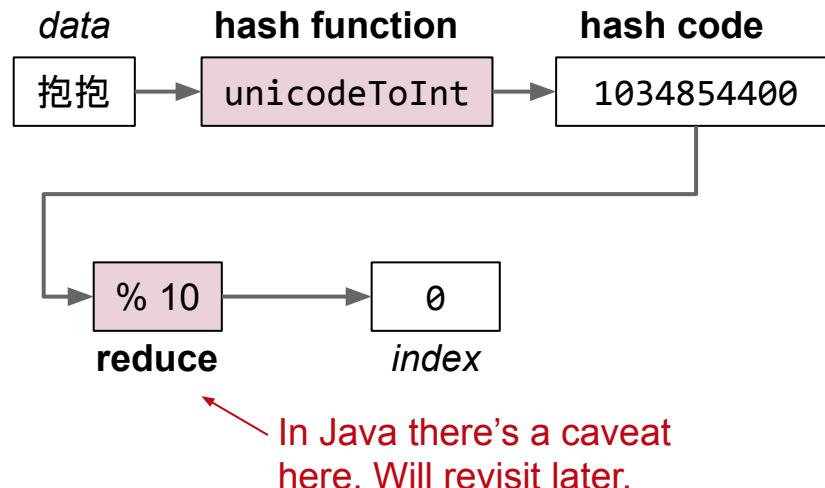
- Put in bucket = hashCode % 10



# The Hash Table

What we've just created here is called a **hash table**.

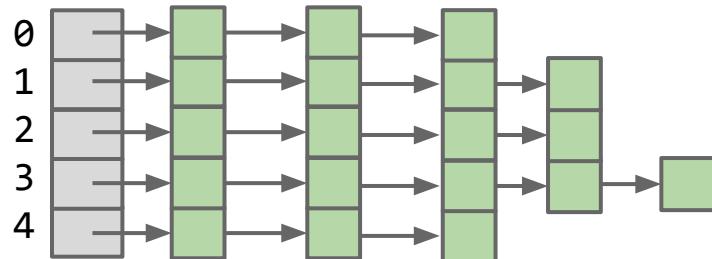
- *Data* is converted by a **hash function** into an integer representation called a **hash code**.
- The **hash code** is then **reduced** to a valid *index*, usually using the modulus operator, e.g.  $2348762878 \% 10 = 8$ .



# Hash Table Performance



# Hash Table Runtime



Worst case runtimes

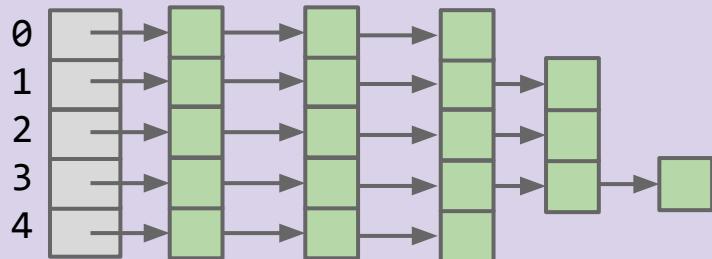
	contains(x)	add(x)
Bushy BSTs	$\Theta(\log N)$	$\Theta(\log N)$
DataIndexedArray	$\Theta(1)$	$\Theta(1)$
Separate Chaining Hash Table	$\Theta(Q)$	$\Theta(Q)$

Q: Length of longest list

The good news: We use way less memory and can now handle arbitrary data.

The bad news: Worst case runtime is now  $\Theta(Q)$ , where Q is the length of the longest list.

# Hash Table Runtime: <http://yellkey.com/somebody>



For the hash table above with 5 buckets, give the order of growth of Q with respect to N.

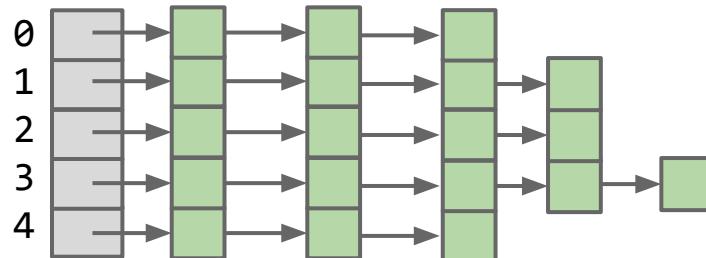
- A. Q is  $\Theta(1)$
- B. Q is  $\Theta(\log N)$
- C. Q is  $\Theta(N)$
- D. Q is  $\Theta(N \log N)$
- E. Q is  $\Theta(N^2)$

## Worst case runtimes

	contains(x)	add(x)
Bushy BSTs	$\Theta(\log N)$	$\Theta(\log N)$
DataIndexedArray	$\Theta(1)$	$\Theta(1)$
Separate Chaining Hash Table	$\Theta(Q)$	$\Theta(Q)$

Q: Length of longest list

# Hash Table Runtime



For the hash table above with 5 buckets, give the order of growth of Q with respect to N.

C.  $Q$  is  $\Theta(N)$

All items evenly distributed.



In the best case, the length of the longest list will be  $N/5$ . In the worst case, it will be  $N$ . In both cases,  $Q(N)$  is  $\Theta(N)$ .

All items in same list.

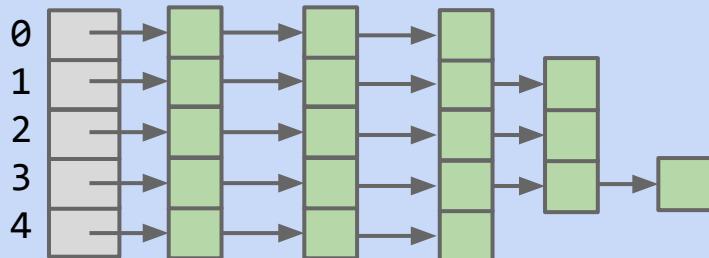


## Worst case runtimes

	contains(x)	add(x)
Bushy BSTs	$\Theta(\log N)$	$\Theta(\log N)$
DataIndexedArray	$\Theta(1)$	$\Theta(1)$
Separate Chaining Hash Table	$\Theta(Q)$	$\Theta(Q)$

Q: Length of longest list

# Improving the Hash Table



Suppose we have:

- A fixed number of buckets  $M$ .
- An increasing number of items  $N$ .

Major problem: Even if items are spread out evenly, lists are of length  $Q = N/M$ .

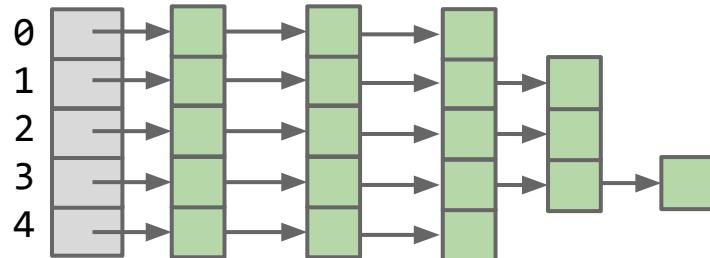
- For  $M = 5$ , that means  $Q = \Theta(N)$ . Results in linear time operations.
- Hard question: How can we improve our design to guarantee that  $N/M$  is  $\Theta(1)$ ?

Worst case runtimes

	contains( $x$ )	add( $x$ )
Bushy BSTs	$\Theta(\log N)$	$\Theta(\log N)$
DataIndexedArray	$\Theta(1)$	$\Theta(1)$
Separate Chaining Hash Table	$\Theta(Q)$	$\Theta(Q)$

$Q$ : Length of longest list

# Hash Table Runtime



Suppose we have:

- An increasing number of buckets M.
- An increasing number of items N.

Major problem: Even if items are spread out evenly, lists are of length  $Q = N/M$ .

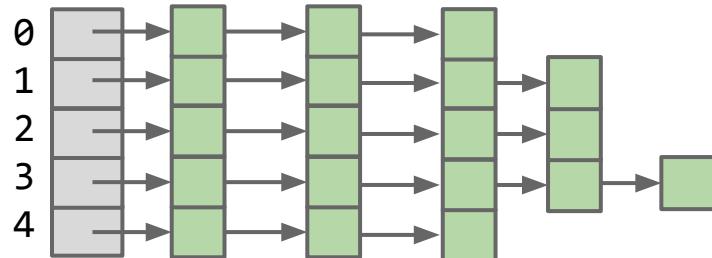
- For  $M = 5$ , that means  $Q = \Theta(N)$ . Results in linear time operations.
- Hard question: How can we improve our design to guarantee that  $N/M$  is  $\Theta(1)$ ?

Worst case runtimes

	contains(x)	add(x)
Bushy BSTs	$\Theta(\log N)$	$\Theta(\log N)$
DataIndexedArray	$\Theta(1)$	$\Theta(1)$
Separate Chaining Hash Table	$\Theta(Q)$	$\Theta(Q)$

Q: Length of longest list

# Hash Table Runtime



Suppose we have:

- An increasing number of buckets M.
- An increasing number of items N.

As long as  $M = \Theta(N)$ , then  $O(N/M) = O(1)$ .

One example strategy: When  $N/M \geq 1.5$ , then double M.

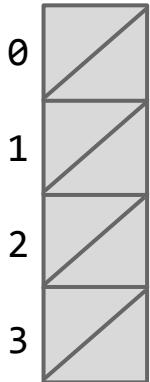
- We often call this process of increasing M “resizing”.
- $N/M$  is often called the “load factor”. It represents how full the hash table is.
- This rule ensures that the average list is never more than 1.5 items long!

# Hash Table Resizing Example

---

When  $N/M \geq 1.5$ , then double M.

$$N = 0 \quad M = 4 \quad N / M = 0$$

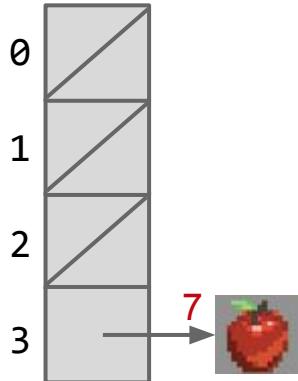


# Hash Table Resizing Example

---

When  $N/M \geq 1.5$ , then double M.

$$N = 1 \quad M = 4 \quad N / M = 0.25$$

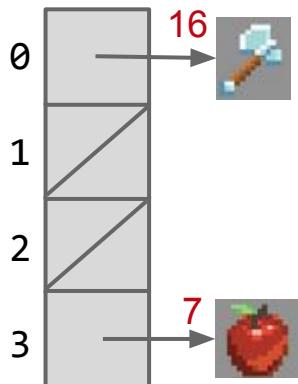


# Hash Table Resizing Example

---

When  $N/M \geq 1.5$ , then double M.

$$N = 2 \quad M = 4 \quad N / M = 0.5$$

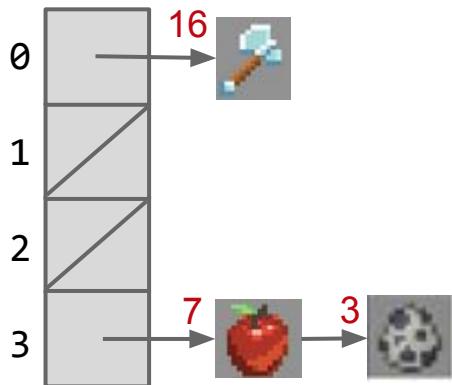


# Hash Table Resizing Example

---

When  $N/M \geq 1.5$ , then double M.

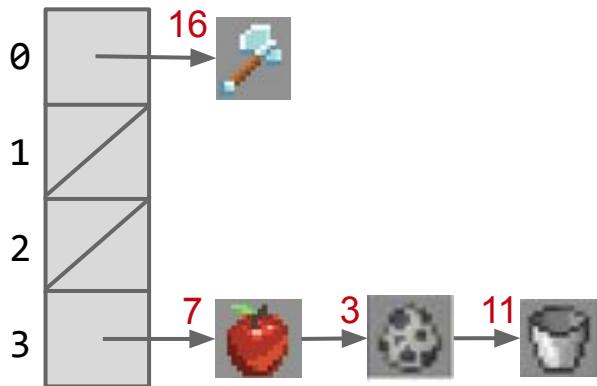
$$N = 3 \quad M = 4 \quad N / M = 0.75$$



# Hash Table Resizing Example

When  $N/M \geq 1.5$ , then double M.

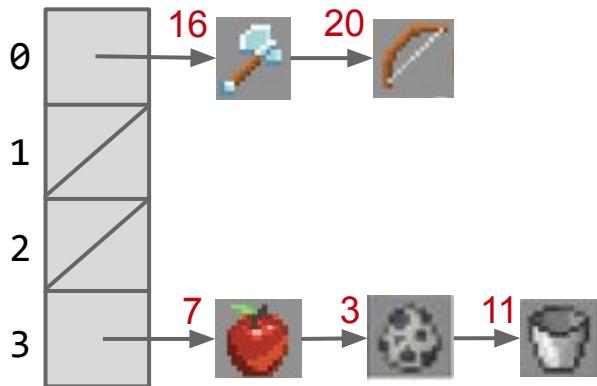
$$N = 4 \quad M = 4 \quad N / M = 1$$



# Hash Table Resizing Example

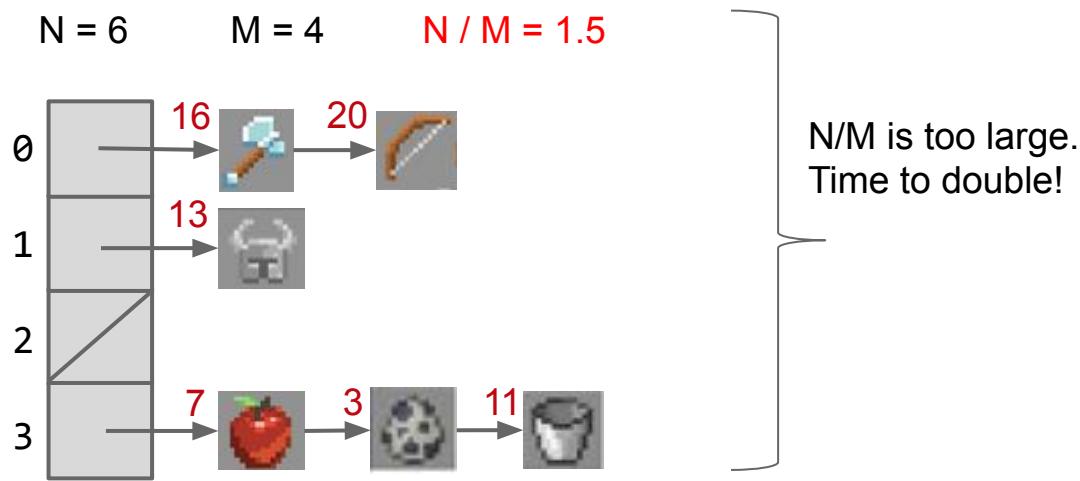
When  $N/M \geq 1.5$ , then double M.

$$N = 5 \quad M = 4 \quad N / M = 1.25$$



# Hash Table Resizing Example

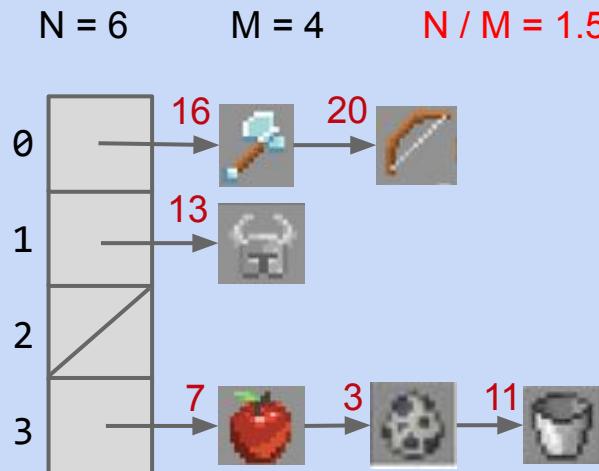
When  $N/M \geq 1.5$ , then double M.



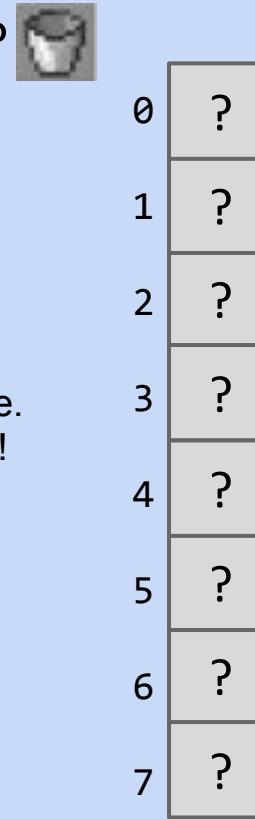
# Hash Table Resizing Example

When  $N/M \geq 1.5$ , then double M.

- Yellkey question: Where will the bucket go?
- Or: Draw the results after doubling M.



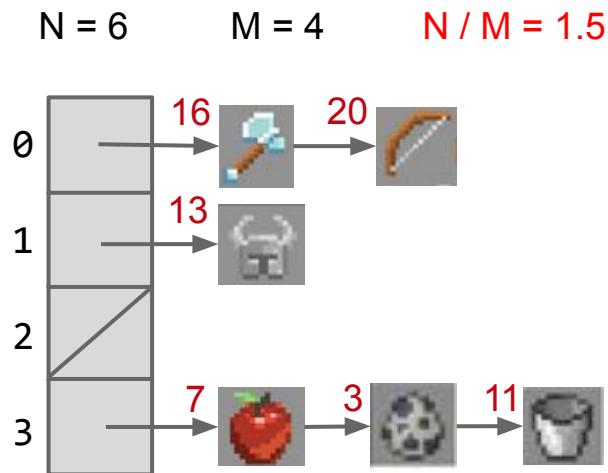
N/M is too large.  
Time to double!



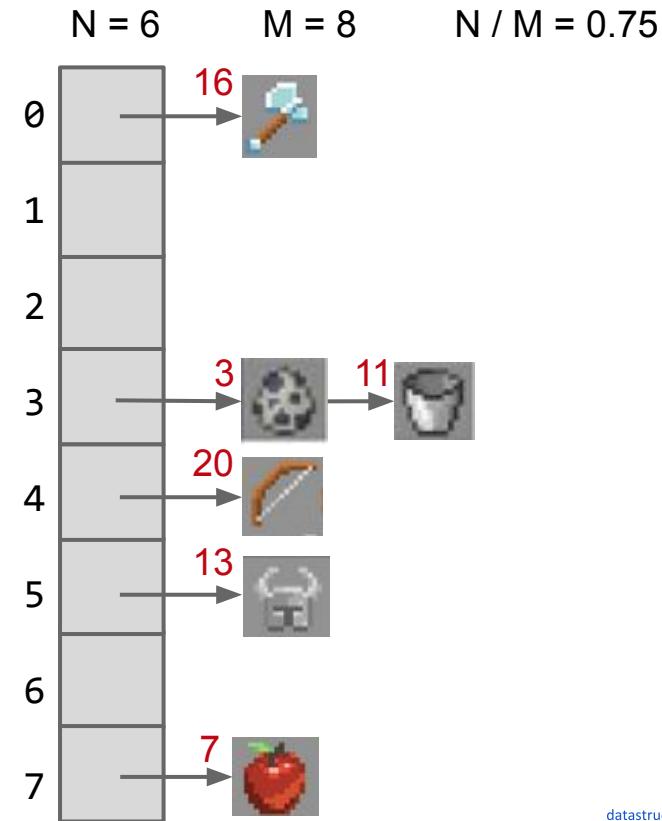
# Hash Table Resizing Example

When  $N/M \geq 1.5$ , then double M.

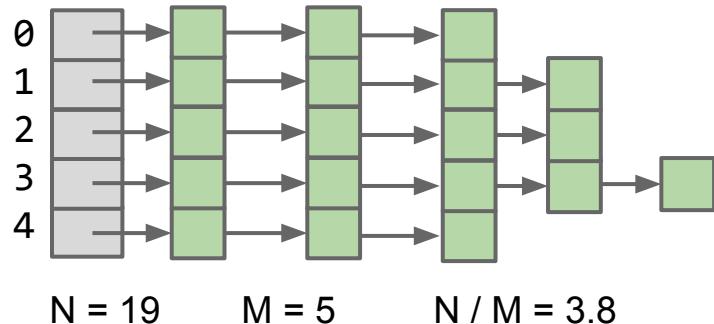
- Draw the results after doubling M.



$N/M$  is too large.  
Time to double!



# Resizing Hash Table Runtime



Suppose we have:

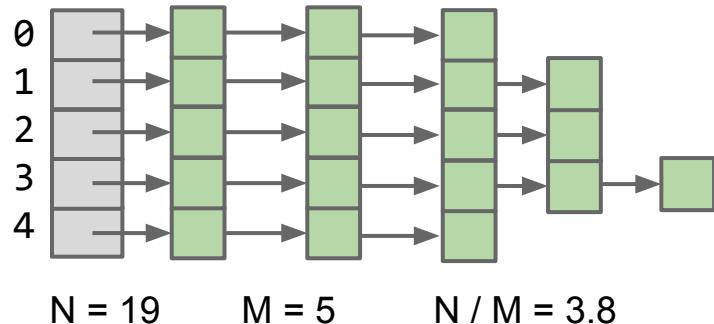
- An increasing number of buckets M.
- An increasing number of items N.

As long as  $M = \Theta(N)$ , then  $O(N/M) = O(1)$ .

*Assuming items are evenly distributed* (as above), lists will be approximately  $N/M$  items long, resulting in  $\Theta(N/M)$  runtimes.

- Our doubling strategy ensures that  $N/M = O(1)$ .
- Thus, worst case runtime for all operations is  $\Theta(N/M) = \Theta(1)$ .
  - ... unless that operation causes a resize.

# Resizing Hash Table Runtime



Suppose we have:

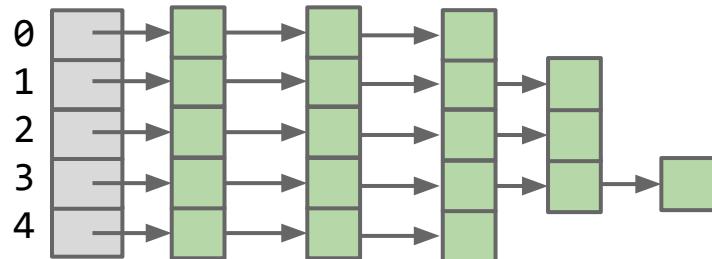
- An increasing number of buckets  $M$ .
- An increasing number of items  $N$ .

As long as  $M = \Theta(N)$ , then  $O(N/M) = O(1)$ .

One important thing to consider is the cost of the resize operation.

- Resizing takes  $\Theta(N)$  time. Have to redistribute all items!
- Most add operations will be  $\Theta(1)$ . Some will be  $\Theta(N)$  time (to resize).
  - Similar to our ALists, as long as we resize by a multiplicative factor, the average runtime will still be  $\Theta(1)$ .
  - Note: We will eventually analyze this in more detail.

# Hash Table Runtime



Because  $Q = \Theta(N)$

Because  $Q = \Theta(1)$

Hash table operations are on average constant time if:

- We double M to ensure constant average bucket length.
- Items are evenly distributed.

## Worst case runtimes

	contains(x)	add(x)
Bushy BSTs	$\Theta(\log N)$	$\Theta(\log N)$
DataIndexedArray	$\Theta(1)$	$\Theta(1)$
Separate Chaining Hash Table With No Resizing	$\Theta(N)$	$\Theta(N)$
... With Resizing	$\Theta(1)^†$	$\Theta(1)^{*†}$

\*: Indicates “on average”.

†: Assuming items are evenly spread.

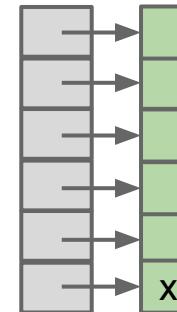
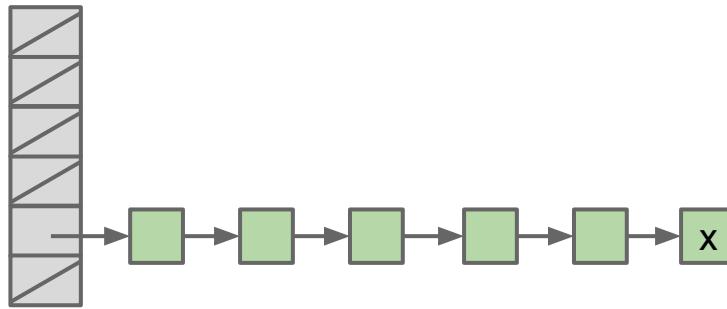
# Regarding Even Distribution

Even distribution of item is critical for good hash table performance.

- Both tables below have load factor of  $N/M = 1$ .
- Left table is much worse!
  - Contains is  $\Theta(N)$  for  $x$ .

Will need to discuss how to ensure even distribution.

- First, let's talk a little bit about how hash tables work in Java.



# Hash Tables in Java



# The Ubiquity of Hash Tables

---

Hash tables are the most popular implementation for sets and maps.

- Great performance in practice.
- Don't require items to be comparable.
- Implementations often relatively simple.
- Python dictionaries are just hash tables in disguise.

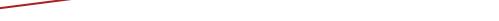
In Java, implemented as `java.util.HashMap` and `java.util.HashSet`.

- How does a `HashMap` know how to compute each object's hash code?
  - Good news: It's not “implements `Hashable`”.
  - Instead, all objects in Java must implement a `.hashCode()` method.

# Objects

---

All classes are hyponyms of Object.

- `String toString()`
- `boolean equals(Object obj)`
- `Class<?> getClass()`
- `int hashCode()` 
- `protected Object clone()`
- `protected void finalize()`
- `void notify()`
- `void notifyAll()`
- `void wait()`
- `void wait(long timeout)`
- `void wait(long timeout, int nanos)`

Default implementation  
simply returns the memory  
address of the object.

# Examples of Real Java HashCodes

We can see that Strings in Java override hashCode, doing something vaguely like what we did earlier.

- Will see the actual hashCode() function later!

```
System.out.println("a".hashCode());
System.out.println("bee".hashCode());
System.out.println("포옹".hashCode());
System.out.println("kamala lifefully".hashCode());
System.out.println("đậu hũ".hashCode());
```

```
jug ~/Dropbox/61b/lec/hashing
$ java JavaHashCodeExamples
    "a"      97
    "bee"    97410
    "포옹"   1732557
"kamala lifefully" 1732557
    "đậu hũ" -2108180664
```

# Using Negative hash codes: <http://yellkey.com/without>

---



Suppose that \_\_\_\_\_'s hash code is -1.

- Philosophically, into which bucket is it most natural to place this item?



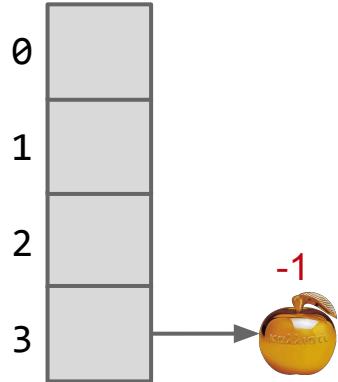
# Using Negative hash codes: <http://yellkey.com/medical>

---



Suppose that \_\_\_\_\_ 's hash code is -1.

- Philosophically, into which bucket is it most natural to place this item?
  - I say 3, since  $-1 \rightarrow 3, 0 \rightarrow 0, 1 \rightarrow 1, 2 \rightarrow 2, 3 \rightarrow 3, 4 \rightarrow 0, \dots$



# Using Negative hash codes in Java



Suppose that \_\_\_\_\_'s hash code is -1.

- Unfortunately,  $-1 \% 4 = -1$ . Will result in index errors!
- Use Math.floorMod instead.

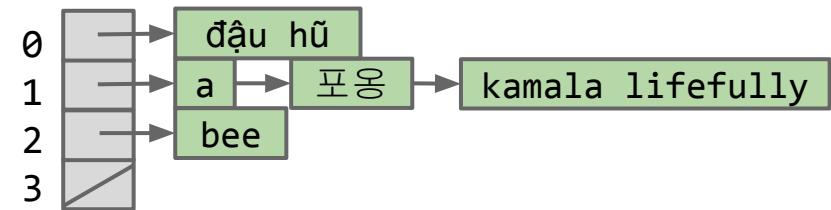
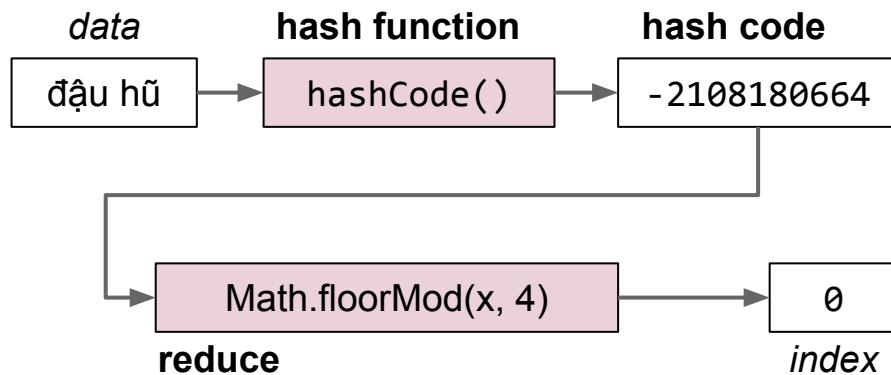
```
0   public class ModTest {  
1     public static void main(String[] args) {  
2       System.out.println(-1 % 4);  
3       System.out.println(Math.floorMod(-1, 4));  
4     }  
5   }
```

```
$ java ModTest  
-1  
3
```

# Hash Tables in Java

Java hash tables:

- *Data* is converted by the **hashCode** method an integer representation called a **hash code**.
- The **hash code** is then **reduced** to a valid *index*, using something like the `floorMod` function, e.g. `Math.floorMod(1732557 % 4) = 8.`



# Two Important Warnings When Using HashMaps/HashSets

---

Warning #1: Never store objects that can change in a HashSet or HashMap!

- If an object's variables changes, then its hashCode changes. May result in items getting lost.

Warning #2: Never override equals without also overriding hashCode.

- Can also lead to items getting lost and generally weird behavior.
- HashMaps and HashSets use equals to determine if an item exists in a particular bucket.
- See HW3 or the study guide problems.



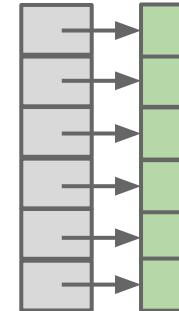
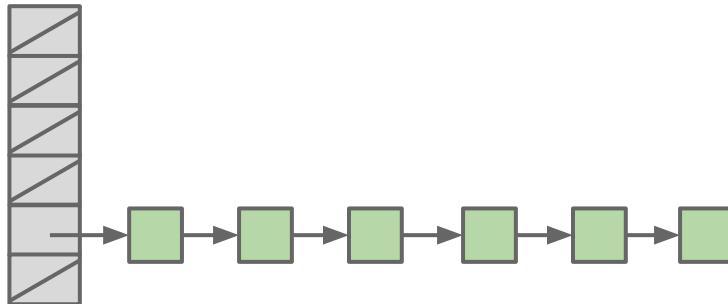
# Good HashCodes (Extra)



# What Makes a good .hashCode()?

Goal: We want hash tables that look like the table on the right.

- Want a hashCode that spreads things out nicely on real data.
  - Example #1: return 0 is a bad hashCode function.
  - Example #2: just returning the first character of a word, e.g. “cat” → 3 was also a bad hash function.
  - Example #3: Adding chars together is bad. “ab” collides with “ba”.
  - Example #4: returning string treated as a base B number can be good.
- Writing a good hashCode() method **can be tricky**.



# Hashbrowns and Hash Codes

---

How do you make hashbrowns?

- Chopping a potato into nice predictable segments? No way!
- Similarly, adding up the characters is not nearly “random” enough.

Can think of multiplying data by powers of some base as ensuring that all the data gets scrambled together into a seemingly random integer.



## Example hashCode Function

The Java 8 hash code for strings. Two major differences from our hash codes:

- Represents strings as a base 31 number.
  - Why such a small base? Real hash codes don't care about uniqueness.
- Stores (caches) calculated hash code so future `hashCode` calls are faster.

```
@Override
public int hashCode() {
    int h = cachedHashCode;
    if (h == 0 && this.length() > 0) {
        for (int i = 0; i < this.length(); i++) {
            h = 31 * h + this.charAt(i);
        }
        cachedHashCode = h;
    }
    return h;
}
```

## Example: Choosing a Base

---

Java's hashCode( ) function for Strings:

- $h(s) = s_0 \times 31^{n-1} + s_1 \times 31^{n-2} + \dots + s_{n-1}$

Our asciiToInt function for Strings:

- $h(s) = s_0 \times 126^{n-1} + s_1 \times 126^{n-2} + \dots + s_{n-1}$

Which is better?

- Might seem like 126 is better. Ignoring overflow, this ensures a unique numerical representation for all ASCII strings.
- ... but overflow is a particularly bad problem for base 126!

## Example: Base 126

---

Major collision problem:

- “geocronite is the best thing on the earth.”.hashCode() yields 634199182.
- “flan is the best thing on the earth.”.hashCode() yields 634199182.
- “treachery is the best thing on the earth.”.hashCode() yields 634199182.
- “Brazil is the best thing on the earth.”.hashCode() yields 634199182.

Any string that ends in the same last 32 characters has the same hash code.

- Why? Because of overflow.
- Basic issue is that  $126^{32} = 126^{33} = 126^{34} = \dots 0$ .
  - Thus upper characters are all multiplied by zero.
  - See CS61C for more.

# Typical Base

---

A typical hash code base is a small prime.

- Why prime?
  - Never even: Avoids the overflow issue on previous slide.
  - Lower chance of resulting hashCode having a bad relationship with the number of buckets: See study guide problems and hw3.
- Why small?
  - Lower cost to compute.

A full treatment of good hash codes is well beyond the scope of our class.

# Hashbrowns and Hash Codes

---

How do you make hashbrowns?

- Chopping a potato into nice predictable segments? No way!

Using a prime base yields better “randomness” than using something like base 126.



## Example: Hashing a Collection

Lists are a lot like strings: Collection of items each with its own hashCode:

```
@Override  
public int hashCode() {  
    int hashCode = 1;  
    for (Object o : this) {  
        hashCode = hashCode * 31; // elevate/smear the current hash code  
        hashCode = hashCode + o.hashCode(); // add new item's hash code  
    }  
    return hashCode;  
}
```

To save time hashing: Look at only first few items.

- Higher chance of collisions but things will still work.

## Example: Hashing a Recursive Data Structure

Computation of the hashCode of a recursive data structure involves recursive computation.

- For example, binary tree hashCode (assuming sentinel leaves):

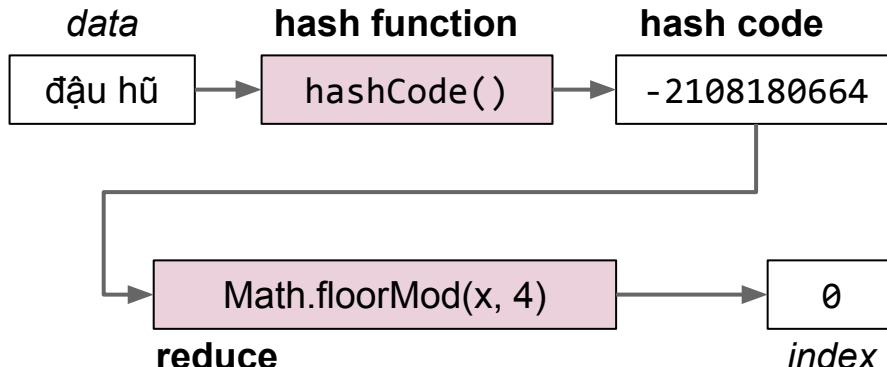
```
@Override  
public int hashCode() {  
    if (this.value == null) {  
        return 0;  
    }  
    return this.value.hashCode() +  
        31 * this.left.hashCode() +  
        31 * 31 * this.right.hashCode();  
}
```

# Summary

# Hash Tables in Java

Hash tables:

- *Data* is converted into a hash code.
- The **hash code** is then **reduced** to a valid *index*.
- *Data* is then stored in a bucket corresponding to that *index*.
- Resize when load factor  $N/M$  exceeds some constant.
- If items are spread out nicely, you get  $\Theta(1)$  average runtime.



	contains(x)	add(x)
Bushy BSTs	$\Theta(\log N)$	$\Theta(\log N)$
Separate Chaining Hash Table With No Resizing	$\Theta(N)$	$\Theta(N)$
... With Resizing	$\Theta(1)^†$	$\Theta(1)^{*†}$

\*: Indicates “on average”.

†: Assuming items are evenly spread.

# Collision Resolution With Linear Probing (Extra)



# Open Addressing: An Alternate Disambiguation Strategy (Extra)

---

An alternate way to handle collisions is to use “open addressing”.

If target bucket is already occupied, use a different bucket, e.g.

- Linear probing: Use next address, and if already occupied, just keep scanning one by one.
  - Demo: <http://goo.gl/o5EDvb>
- Quadratic probing: Use next address, and if already occupied, try looking 4 ahead, then 9 ahead, then 16 ahead, ...
- Many more possibilities. See the optional reading for today (or CS170) for a more detailed look.

In 61B, we'll settle for separate chaining.

## Citations

---

<http://www.nydailynews.com/news/national/couple-calls-911-forgotten-mcdonalds-hash-browns-article-1.1543096>

[http://en.wikipedia.org/wiki/Pigeonhole\\_principle#mediaviewer/File:TooManyPigeons.jpg](http://en.wikipedia.org/wiki/Pigeonhole_principle#mediaviewer/File:TooManyPigeons.jpg)

[https://cookingplanit.com/public/uploads/inventory/hashbrown\\_1366322674.jpg](https://cookingplanit.com/public/uploads/inventory/hashbrown_1366322674.jpg)